

Assignment 1

Course: Deep Learning

Instructor: Dr. Ali Irtaza

Department of Electrical Engineering

Institute of Space Technology

Submitted by:

Zawar Khan

September 22, 2026

Q3. Consider two sets of outputs

Set A: outcomes with *minimum to low* entropy

Set B: outcomes with *medium to high* entropy

(a) Which set is more meaningful for decreasing unpredictable behaviour?

Shannon entropy (analytical)

Shannon's information theory defines the unpredictability of a discrete random variable [?]. The Shannon measure of unpredictability is the **Shannon entropy**. For a single outcome i with probability p_i , Shannon defined the *surprise* (self-information) as

$$\text{information}(i) = \log_b \left(\frac{1}{p_i} \right) = -\log_b(p_i). \quad (1)$$

Interpretation of (??):

- If p_i is close to 1 (a near-certain event), then $\log(1/p_i)$ is close to 0 — almost no information/surprise.
- If p_i is close to 0 (a rare event), then $\log(1/p_i)$ is large — a large amount of information/surprise.

For *many* outcomes, the entropy of the full distribution is the **probability-weighted sum** of these surprises:

$$H(X) = \sum_i p_i \text{information}(i) = \sum_i p_i (-\log_b p_i) = -\sum_i p_i \log_b p_i, \quad (2)$$

where p_i is the probability of outcome i and the base b chooses the unit (bits when $b = 2$). Equation (??) is the general multi-outcome formula used for Set A and Set B below.

Special case (two outcomes). For a Bernoulli system with probabilities p and $1 - p$, (??) reduces to the **binary entropy**

$$H(p) = -p \log_2 p - (1 - p) \log_2 (1 - p). \quad (3)$$

$H(p)$ is *low* when p is near 0 or near 1 (highly predictable outcomes) and reaches its *maximum* at $p = 0.5$ (maximum randomness). Thus “minimum to low entropy” corresponds to probability mass concentrated on nearly certain outcomes, while “medium to high entropy” corresponds to more uniformly uncertain outcomes.

Figures

Figure ?? shows the binary-entropy special case $H(p)$. Entropy is smallest near the extremes of p .




Figure 1: Binary Shannon entropy $H(p)$ versus probability p (special case of two outcomes). Low entropy occurs for very small or very large p .

To connect (??) to Set A and Set B, we form two sample sets whose outcomes are integers from 1 to 10 (size $N = 100$):

- **Set A:** draws are random, but a few numbers appear much more often (peaked frequencies $\Rightarrow$ low entropy).
- **Set B:** draws are nearly uniform, so the same number rarely dominates (spread frequencies $\Rightarrow$ high entropy).

Empirical probabilities are $p_i = n_i/N$. Importantly:

- The **top histograms** show only *counts* (how often each outcome appears). They are *not* an entropy curve.
- **Entropy** is a *single number* for each whole set, $H = -\sum_i p_i \log_2 p_i$, shown in the **bottom panel** of Figure ??.




Figure 2: Top: outcome frequencies for Set A and Set B (values 1–10). Bottom: Shannon entropy of each set (one bar per set). The histograms describe the data; the bottom panel is where entropy is compared.

Conclusion

Entropy is a direct measure of randomness / unpredictability. To *reduce* unpredictable behaviour we need outcomes that are as predictable as possible — concentrated probability mass and **low Shannon entropy**. From (??), (??), and Figures ??–??. that regime is exactly **Set A**. Set B corresponds to greater uncertainty and is less suitable if the design goal is to reduce randomness.

Note. Full plotting scripts and libraries are in the course repository [?].

Python Implementation

```
1  import numpy as np
2
3  def shannon_entropy(probs, base=2.0):
4      """ $H = -\sum p_i \log_b(p_i)$  (weighted sum over many
5          outcomes)."""
6      p = np.asarray(probs, dtype=float)
7      p = p / p.sum()
8      nz = p > 0
9      return float(-(p[nz] * (np.log(p[nz]) / np.log(base))).sum
10                     ())
11
12 def binary_entropy(p, base=2.0):
13     """Special case:  $H(p)$  for Bernoulli( $p$ ); low near  $p=0$  or  $p=1$ ."""
14     p = np.clip(np.asarray(p, dtype=float), 0.0, 1.0)
15     h = np.zeros_like(p, dtype=float)
16     mid = (p > 0) & (p < 1)
17     pm = p[mid]
18     h[mid] = -(pm * np.log(pm) + (1 - pm) * np.log(1 - pm)) /
19               np.log(base)
20     return h
```

Listing 1: Shannon entropy and binary entropy

Q4. Normal distribution and maximum entropy

If we know the mean μ and standard deviation σ , is it true that the Normal (Gaussian) distribution has the maximum entropy?

Differential entropy

For a continuous random variable with density $p(x)$, Shannon's discrete sum becomes an integral — the **differential entropy**

$$H(X) = - \int_{-\infty}^{\infty} p(x) \log p(x) dx \quad (4)$$

(with base 2 if the answer is wanted in bits). This is the continuous analogue of (??). Standard results in information theory show that, among all continuous densities that share a fixed mean μ and fixed variance σ^2 , the Gaussian density uniquely maximises $H(X)$ [?].

Entropy of a Gaussian

For $X \sim \mathcal{N}(\mu, \sigma^2)$ the integral (??) evaluates in closed form to

$$H(X) = \frac{1}{2} \log_2(2\pi e \sigma^2) \quad (\text{bits}). \quad (5)$$

Two observations follow immediately from (??):

- $H(X)$ depends only on the variance σ^2 , *not* on the mean μ . Entropy measures unpredictability (spread), not where the bell curve is centred. Shifting the curve left or right does not change how uncertain the outcome is.
- As $\sigma \rightarrow 0$ (a very narrow spike) the entropy becomes very small: you are almost certain of the value. As σ grows (a very wide curve) the entropy grows: values can land almost anywhere, so uncertainty is high.

Maximum-entropy principle (why the Gaussian?)

The question is not “does a Gaussian have high or low entropy in absolute terms?” It is: *given only* the constraints

1. $\int p(x) dx = 1$ (a valid density),
2. $\mathbb{E}[X] = \mu$ and $\text{Var}(X) = \sigma^2$ (known mean and spread),

which density p makes $H(X)$ as large as possible?

The answer is the Normal density. Equivalently: among all continuous distributions that share the same μ and σ^2 , the Gaussian is the one with the **highest** differential entropy. Any other shape with the same mean and variance (skewed, multimodal, truncated, ...) secretly encodes *extra* structure beyond what μ and σ^2 alone justify, and that extra structure *reduces* entropy. Choosing the Gaussian is therefore the “least additional-assumption” choice consistent with those two facts — the **maximum entropy principle**, which is why the Normal appears so often as the default model when only mean and variance are known [?].

Clarifying the Empirical Rule (68–95–99.7)

For a Gaussian, about 68% of the probability mass lies within 1σ of μ , 95% within 2σ , and 99.7% within 3σ . That rule describes *where* the Gaussian places its mass given σ ; it does *not* mean “entropy must be small” or that the Gaussian is more predictable than every other distribution.

The comparison that matters for this question is across distributions that all have the *same* σ^2 :

- Every such distribution has the same variance constraint, so they are equally “spread” in the second-moment sense.
- The Empirical Rule is a property of the Gaussian shape under that constraint; other shapes can put mass differently (e.g. sharper peaks or hard cut-offs) and then have *lower* entropy.
- Knowing “most samples fall within a few σ ” for a Gaussian is already built into (??); it does not contradict maximality. Maximality says: with only μ and σ^2 known, you cannot claim a more structured (lower entropy) law without adding assumptions you were not given.

In short: variance σ^2 sets the entropy *scale* for a Gaussian via (??); the maximum-entropy principle says that, *under the same* μ, σ^2 , no competing continuous density beats the Gaussian.

Illustration (same μ, σ^2 , different shapes)

Figure ?? fixes $\mu = 0$ and $\sigma = 1$ and plots three densities that all have this mean and variance: Gaussian, Laplace, and Uniform. The left panel shows the shapes; the right panel shows their differential entropies. The Gaussian bar is the tallest — that is the content of the maximum-entropy claim.




Figure 3: Left: three PDFs with identical μ and σ^2 . Right: differential entropy of each. The Gaussian has the largest H ; the other shapes add extra structure (sharper peak or hard cut-offs) and therefore lower entropy.

Answer

Yes. If the mean μ and standard deviation σ (equivalently variance σ^2) are fixed, the Normal distribution has the maximum differential entropy among all continuous distributions with those moments. Its entropy is $H = \frac{1}{2} \log_2(2\pi e\sigma^2)$ and depends on σ , not on μ .

Q5. Find and plot the derivatives of (a) ReLU, (b) Sigmoid, and (c) Hyperbolic Tangent functions.

(a) ReLU

Analytical Form

The Rectified Linear Unit (ReLU) is defined as

$$\text{ReLU}(x) = \max(0, x) = \begin{cases} 0, & x < 0, \\ x, & x \geq 0. \end{cases} \quad (6)$$

Differentiating piecewise (taking the subgradient at $x = 0$ to be 0, as is common in practice) yields

$$\text{ReLU}'(x) = \begin{cases} 0, & x \leq 0, \\ 1, & x > 0. \end{cases} \quad (7)$$

Figure 1 shows the ReLU activation function and its derivative.

Plot

The derivative in python code is simple finite check in the limits given while for the derivative at Zero its taken as 0.




Figure 4: ReLU activation (top) and its derivative (bottom).

Python Implementation

For the computation of derivatives we have used the NumPy library.

```
1 import numpy as np
2 def relu(x):
3     return np.maximum(0.0, x)
4
5 def relu_derivative(x):
6     return (x > 0).astype(float)
```

Listing 2: ReLU and its derivative

(b) Sigmoid

Analytical Form

The logistic sigmoid is

$$\sigma(x) = \frac{1}{1 + e^{-x}}. \quad (8)$$

Differentiating with respect to x :

$$\sigma'(x) = \frac{e^{-x}}{(1 + e^{-x})^2} = \frac{1}{1 + e^{-x}} \cdot \frac{e^{-x}}{1 + e^{-x}} = \sigma(x) (1 - \sigma(x)). \quad (9)$$

Thus

$$\sigma'(x) = \sigma(x) (1 - \sigma(x)). \quad (10)$$

Figure

Figure 2 shows the Sigmoid activation function and its derivative. Chain rule has been used to find the derivative.




Figure 5: Sigmoid activation (top) and its derivative (bottom).

Python Implementation

```
1 import numpy as np
2
3 def sigmoid(x):
4     return 1.0 / (1.0 + np.exp(-x))
5
6 def sigmoid_derivative(x):
7     s = sigmoid(x)
8     return s * (1.0 - s)
```

Listing 3: Sigmoid and its derivative

(c) Hyperbolic Tangent

Analytical Form

The hyperbolic tangent is

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = \frac{\sinh(x)}{\cosh(x)}. \quad (11)$$

Its derivative follows from the quotient rule (or the identity $\frac{d}{dx} \tanh x = \text{sech}^2 x$):

$$\tanh'(x) = 1 - \tanh^2(x) = \text{sech}^2(x). \quad (12)$$

Figure

Figure 3 shows the Hyperbolic Tangent activation function and its derivative. Chain rule has been used to find the derivative.




Figure 6: Hyperbolic tangent activation (top) and its derivative (bottom).

Python Implementation

```
1 import numpy as np
2
3 def tanh_fn(x):
4     return np.tanh(x)
5
6 def tanh_derivative(x):
7     t = np.tanh(x)
8     return 1.0 - t**2
```

Listing 4: Hyperbolic tangent and its derivative

References

- [1] Q. Peter He, *Information Entropy based Indices for Variable Selection Performance Assessment*, in Mario R. Eden, Marianthi G. Ierapetritou, Gavin P. Towler (Eds.), *Computer Aided Chemical Engineering*, Elsevier, Volume 44, 2018, pp. 2281–2286, ISSN 1570-7946, ISBN 9780444642417, <https://doi.org/10.1016/B978-0-444-64241-7.50375-X>.
- [2] Z. Khan, *Deep Learning Course* (source code and scripts), GitHub repository, https://github.com/ZawarKhan97/Deep_Learning_Course.
- [3] R. E. Blahut, *Information Theory and Coding*, in Wendy M. Middleton, Mac E. Van Valkenburg (Eds.), *Reference Data for Engineers* (Ninth Edition), Newnes, 2002, pp. 25-1–25-31, ISBN 9780750672917, <https://doi.org/10.1016/B978-075067291-7/50027-3>.